

Tutorial: Actor Critic Implementation

In [1]:

```
1  # Import required libraries
2
3  import argparse
4  import gymnasium as gym
5  import numpy as np
6  from collections import namedtuple
7
8  import torch
9  import torch.nn as nn
10 import torch.nn.functional as F
11 import torch.optim as optim
12 from torch.distributions import Categorical
```

In [2]:

```
1  # Set constants for training # DO NOT CHANGE
2  seed = 543
3  log_interval = 10
4  gamma = 0.99
5
6  env = gym.make('CartPole-v1')
7  env.reset(seed=seed)
8  torch.manual_seed(seed)
9
10 SavedAction = namedtuple('SavedAction', ['log_prob', 'value'])
```

In [3]:

```
1
2  class Policy(nn.Module):
3      """
4      implements both actor and critic in one model
5      """
6
7      def __init__(self):
8          super(Policy, self).__init__()
9          self.affine1 = nn.Linear(4, 128)
10
11         # actor's layer
12         self.action_head = nn.Linear(128, 2)
13
14         # critic's layer
15         self.value_head = nn.Linear(128, 1)
16
17         # action & reward buffer
18         self.saved_actions = []
19         self.rewards = []
20
21     def forward(self, x):
22         """
23         forward of both actor and critic
24         """
25         x = F.relu(self.affine1(x))
26
27         # actor: choses action to take from state s_t
28         # by returning probability of each action
29         action_prob = F.softmax(self.action_head(x), dim=-1)
30
31         # critic: evaluates being in the state s_t
32         state_values = self.value_head(x)
33
34         # return values for both actor and critic as a tuple of 2 values:
35         # 1. a list with the probability of each action over the action space
36         # 2. the value from state s_t
37         return action_prob, state_values
```


In [4]:

```

1  # DO NOT Modify Training Code
2
3  def select_action(state):
4      state = torch.from_numpy(state).float()
5      probs, state_value = model(state)
6
7      # create a categorical distribution over the list of probabilities of actions
8      m = Categorical(probs)
9
10     # and sample an action using the distribution
11     action = m.sample()
12
13     # save to action buffer
14     model.saved_actions.append(SavedAction(m.log_prob(action), state_value))
15
16     # the action to take (left or right)
17     return action.item()
18
19
20 def finish_episode():
21     """
22     Training code. Calculates actor and critic loss and performs backprop.
23     """
24     R = 0
25     saved_actions = model.saved_actions
26     policy_losses = [] # list to save actor (policy) loss
27     value_losses = [] # list to save critic (value) loss
28     returns = [] # list to save the true values
29
30     # calculate the true value using rewards returned from the environment
31     for r in model.rewards[::-1]:
32         # calculate the discounted value
33         R = r + gamma * R
34         returns.insert(0, R)
35
36     returns = torch.tensor(returns)
37     returns = (returns - returns.mean()) / (returns.std() + eps)
38
39     for (log_prob, value), R in zip(saved_actions, returns):
40         advantage = R - value.item()
41
42         # calculate actor (policy) loss
43         policy_losses.append(-log_prob * advantage)
44
45         # calculate critic (value) loss using L1 smooth loss
46         value_losses.append(F.smooth_l1_loss(value, torch.tensor([R])))
47
48     # reset gradients
49     optimizer.zero_grad()
50
51     # sum up all the values of policy_losses and value_losses
52     loss = torch.stack(policy_losses).sum() + torch.stack(value_losses).sum()
53
54     # perform backprop
55     loss.backward()
56     optimizer.step()
57
58     # reset rewards and action buffer
59     del model.rewards[:]
60     del model.saved_actions[:]
61
62
63 def train():
64     running_reward = 10
65
66     # run infinitely many episodes
67     for i_episode in range(2000):
68
69         # reset environment and episode reward
70         state = env.reset()[0]
71         ep_reward = 0
72
73         # for each episode, only run 9999 steps so that we don't
74         # infinite loop while learning
75         for t in range(1, 10000):
76
77             # select action from policy
78             action = select_action(state)
79
80             # take the action
81             state, reward, done, truncated, _ = env.step(action)
82
83             model.rewards.append(reward)
84             ep_reward += reward
85             if done:
86                 break
87
88             # update cumulative reward

```

```
89     running_reward = 0.05 * ep_reward + (1 - 0.05) * running_reward
90
91     # perform backprop
92     finish_episode()
93
94     # Log results
95     if i_episode % log_interval == 0:
96         print('Episode {} \t Last reward: {:.2f} \t Average reward: {:.2f}'.format(
97             i_episode, ep_reward, running_reward))
98
99     # check if we have "solved" the cart pole problem
100    if running_reward > env.spec.reward_threshold:
101        print("Solved! Running reward is now {} and "
102              "the last episode runs to {} time steps!".format(running_reward, t))
103        break
```

```
In [14]: 1 %%time
2         # Trail 1
3         model = Policy()
4         optimizer = optim.Adam(model.parameters(), lr=1.5e-2)
5         eps = np.finfo(np.float32).eps.item()
6         train()
```

Episode 0 Last reward: 28.00 Average reward: 10.90
Episode 10 Last reward: 63.00 Average reward: 19.10
Episode 20 Last reward: 322.00 Average reward: 59.17
Episode 30 Last reward: 82.00 Average reward: 110.03
Episode 40 Last reward: 98.00 Average reward: 101.12
Episode 50 Last reward: 92.00 Average reward: 92.00
Episode 60 Last reward: 256.00 Average reward: 176.70
Episode 70 Last reward: 268.00 Average reward: 188.68
Solved! Running reward is now 670.8592141216325 and the last episode runs to 9319 time steps!
CPU times: user 20.2 s, sys: 35.1 ms, total: 20.3 s
Wall time: 20.9 s

```
In [15]: 1 %%time
2         # Trail 2
3         model = Policy()
4         optimizer = optim.Adam(model.parameters(), lr=1.5e-2)
5         eps = np.finfo(np.float32).eps.item()
6         train()
```

Episode 0 Last reward: 14.00 Average reward: 10.20
Episode 10 Last reward: 11.00 Average reward: 12.16
Episode 20 Last reward: 10.00 Average reward: 14.00
Episode 30 Last reward: 15.00 Average reward: 15.94
Episode 40 Last reward: 34.00 Average reward: 19.84
Episode 50 Last reward: 81.00 Average reward: 38.05
Episode 60 Last reward: 197.00 Average reward: 63.05
Episode 70 Last reward: 200.00 Average reward: 109.05
Episode 80 Last reward: 41.00 Average reward: 87.67
Episode 90 Last reward: 242.00 Average reward: 89.43
Episode 100 Last reward: 97.00 Average reward: 170.49
Episode 110 Last reward: 311.00 Average reward: 212.84
Episode 120 Last reward: 437.00 Average reward: 261.66
Episode 130 Last reward: 219.00 Average reward: 316.11
Episode 140 Last reward: 44.00 Average reward: 249.87
Episode 150 Last reward: 106.00 Average reward: 197.43
Episode 160 Last reward: 213.00 Average reward: 219.95
Solved! Running reward is now 538.783586169595 and the last episode runs to 6095 time steps!
CPU times: user 29 s, sys: 48.9 ms, total: 29 s
Wall time: 29.1 s

In [17]:

```
1 %%time
2 # Trail 3
3 model = Policy()
4 optimizer = optim.Adam(model.parameters(), lr=1.5e-2)
5 eps = np.finfo(np.float32).eps.item()
6 train()

Episode 0      Last reward: 17.00      Average reward: 10.35
Episode 10     Last reward: 38.00      Average reward: 16.29
Episode 20     Last reward: 104.00     Average reward: 34.82
Episode 30     Last reward: 28.00      Average reward: 61.61
Episode 40     Last reward: 97.00      Average reward: 71.99
Episode 50     Last reward: 87.00      Average reward: 77.92
Episode 60     Last reward: 23.00      Average reward: 76.47
Episode 70     Last reward: 28.00      Average reward: 66.68
Episode 80     Last reward: 158.00     Average reward: 77.86
Episode 90     Last reward: 59.00      Average reward: 90.38
Episode 100    Last reward: 99.00      Average reward: 95.70
Episode 110    Last reward: 529.00     Average reward: 135.16
Episode 120    Last reward: 312.00     Average reward: 178.71
Episode 130    Last reward: 619.00     Average reward: 251.69
Episode 140    Last reward: 926.00     Average reward: 300.43
Solved! Running reward is now 554.9558954068989 and the last episode runs to 5285 time steps!
CPU times: user 24.3 s, sys: 107 ms, total: 24.4 s
Wall time: 24.7 s
```

TODO: Write a policy class similar to the above, without using shared features for the actor and critic and compare their performance.

In [8]:

```
1
2 # TODO: Write a policy class similar to the above, without using shared features for the actor and critic and compar
3 # performance.
4
5 class UnsharedPolicy(nn.Module):
6     def __init__(self):
7         super(UnsharedPolicy, self).__init__()
8         # TODO: Fill in.
9         hidden_size = 128
10        # Actor network
11        self.actor_affine1 = nn.Linear(4, hidden_size)
12        self.action_head = nn.Linear(hidden_size, 2)
13
14        # Critic network
15        self.critic_affine1 = nn.Linear(4, hidden_size)
16        self.value_head = nn.Linear(hidden_size, 1)
17
18        self.saved_actions = []
19        self.rewards = []
20
21    def forward(self, x):
22        # TODO: Fill in. For your networks, use the same hidden_size for the layers as the previous policy, that is
23        # Actor forward pass
24        actor_x = F.relu(self.actor_affine1(x))
25        action_prob = F.softmax(self.action_head(actor_x), dim=-1)
26
27        # Critic forward pass
28        critic_x = F.relu(self.critic_affine1(x))
29        state_values = self.value_head(critic_x)
30        # return values for both actor and critic as a tuple of 2 values:
31        # 1. A list with the probability of each action over the action space
32        # 2. The value from state s_t
33        return action_prob, state_values
34
35
```

```
In [9]: 1 %%time
2 # Trail 1
3 model = UnsharedPolicy()
4 optimizer = optim.Adam(model.parameters(), lr=1.5e-2)
5 eps = np.finfo(np.float32).eps.item()
6 train()
```

Episode 0	Last reward: 12.00	Average reward: 10.10
Episode 10	Last reward: 72.00	Average reward: 21.01
Episode 20	Last reward: 78.00	Average reward: 37.09
Episode 30	Last reward: 32.00	Average reward: 48.21
Episode 40	Last reward: 131.00	Average reward: 91.93
Episode 50	Last reward: 80.00	Average reward: 87.19
Episode 60	Last reward: 105.00	Average reward: 83.80
Episode 70	Last reward: 438.00	Average reward: 134.90
Episode 80	Last reward: 661.00	Average reward: 359.26
Episode 90	Last reward: 215.00	Average reward: 326.31
Episode 100	Last reward: 195.00	Average reward: 289.66
Episode 110	Last reward: 189.00	Average reward: 264.00
Episode 120	Last reward: 108.00	Average reward: 233.10
Episode 130	Last reward: 39.00	Average reward: 182.92
Episode 140	Last reward: 103.00	Average reward: 139.45
Episode 150	Last reward: 104.00	Average reward: 124.55
Episode 160	Last reward: 88.00	Average reward: 111.15
Episode 170	Last reward: 123.00	Average reward: 110.90
Episode 180	Last reward: 154.00	Average reward: 123.97
Episode 190	Last reward: 212.00	Average reward: 149.61
Episode 200	Last reward: 377.00	Average reward: 208.81
Episode 210	Last reward: 406.00	Average reward: 294.45

Solved! Running reward is now 844.579265464063 and the last episode runs to 9999 time steps!

CPU times: user 52.2 s, sys: 148 ms, total: 52.3 s

Wall time: 52.5 s

```
In [11]: 1 %%time
2 # Trail 2
3 model = UnsharedPolicy()
4 optimizer = optim.Adam(model.parameters(), lr=1.5e-2)
5 eps = np.finfo(np.float32).eps.item()
6 train()
```

Episode 0	Last reward: 10.00	Average reward: 10.00
Episode 10	Last reward: 9.00	Average reward: 10.57
Episode 20	Last reward: 12.00	Average reward: 11.45
Episode 30	Last reward: 10.00	Average reward: 11.09
Episode 40	Last reward: 9.00	Average reward: 11.12
Episode 50	Last reward: 11.00	Average reward: 12.45
Episode 60	Last reward: 46.00	Average reward: 17.77
Episode 70	Last reward: 34.00	Average reward: 27.80
Episode 80	Last reward: 77.00	Average reward: 47.31
Episode 90	Last reward: 54.00	Average reward: 49.86
Episode 100	Last reward: 75.00	Average reward: 59.78
Episode 110	Last reward: 101.00	Average reward: 69.92
Episode 120	Last reward: 67.00	Average reward: 73.71
Episode 130	Last reward: 74.00	Average reward: 80.61
Episode 140	Last reward: 40.00	Average reward: 74.51
Episode 150	Last reward: 44.00	Average reward: 67.89
Episode 160	Last reward: 46.00	Average reward: 61.33
Episode 170	Last reward: 106.00	Average reward: 68.41
Episode 180	Last reward: 181.00	Average reward: 92.58
Episode 190	Last reward: 120.00	Average reward: 98.85
Episode 200	Last reward: 133.00	Average reward: 105.84
Episode 210	Last reward: 197.00	Average reward: 135.52
Episode 220	Last reward: 131.00	Average reward: 127.97
Episode 230	Last reward: 152.00	Average reward: 130.90

Solved! Running reward is now 669.9642084534141 and the last episode runs to 9999 time steps!

CPU times: user 31 s, sys: 62.5 ms, total: 31.1 s

Wall time: 33.9 s

In [13]:

1

2

3

4

5

6

%%time

Trail 3

model = UnsharedPolicy()

optimizer = optim.Adam(model.parameters(), lr=1.5e-2)

eps = np.finfo(np.float32).eps.item()

train()

Episode 0	Last reward: 42.00	Average reward: 11.60
Episode 10	Last reward: 10.00	Average reward: 11.06
Episode 20	Last reward: 10.00	Average reward: 10.51
Episode 30	Last reward: 10.00	Average reward: 10.33
Episode 40	Last reward: 9.00	Average reward: 9.96
Episode 50	Last reward: 9.00	Average reward: 9.69
Episode 60	Last reward: 10.00	Average reward: 9.59
Episode 70	Last reward: 11.00	Average reward: 9.73
Episode 80	Last reward: 11.00	Average reward: 9.56
Episode 90	Last reward: 11.00	Average reward: 9.60
Episode 100	Last reward: 12.00	Average reward: 10.66
Episode 110	Last reward: 17.00	Average reward: 13.32
Episode 120	Last reward: 15.00	Average reward: 16.55
Episode 130	Last reward: 49.00	Average reward: 29.21
Episode 140	Last reward: 10.00	Average reward: 24.73
Episode 150	Last reward: 19.00	Average reward: 24.10
Episode 160	Last reward: 104.00	Average reward: 40.09
Episode 170	Last reward: 48.00	Average reward: 43.23
Episode 180	Last reward: 50.00	Average reward: 50.56
Episode 190	Last reward: 48.00	Average reward: 50.56
Episode 200	Last reward: 189.00	Average reward: 65.73
Episode 210	Last reward: 106.00	Average reward: 144.54
Episode 220	Last reward: 93.00	Average reward: 124.08
Episode 230	Last reward: 78.00	Average reward: 106.24
Episode 240	Last reward: 75.00	Average reward: 91.86
Episode 250	Last reward: 78.00	Average reward: 85.44
Episode 260	Last reward: 106.00	Average reward: 82.58
Episode 270	Last reward: 119.00	Average reward: 94.06
Episode 280	Last reward: 115.00	Average reward: 103.90
Episode 290	Last reward: 147.00	Average reward: 110.33
Episode 300	Last reward: 118.00	Average reward: 114.31
Episode 310	Last reward: 113.00	Average reward: 120.09
Episode 320	Last reward: 93.00	Average reward: 115.12
Episode 330	Last reward: 191.00	Average reward: 127.87

Solved! Running reward is now 698.7663821435183 and the last episode runs to 9999 time steps!

CPU times: user 35.7 s, sys: 75.8 ms, total: 35.8 s

Wall time: 36 s

By Running the experiments for 3 consecutive trails, It is observed that the Actor Critic model with shared features perform better than the Actor critic model with unshared features .

The model with shared features learns a common feature representation by which the actor and critic update their parameters. Due to this there is stable updates and faster learning. The trails conducted also indicate the same, where model with shared features learns to solve the environment in lesser number of epochs compared to model with unshared features.

The parameters:

1. CPU time: Total time spent by CPU in executing the task.
- For Model with Shared features:

▪ Avg CPU time over 3 trails: 24.57s

• For Model with Unshared features:

▪ Avg CPU time over 3 trails: 39.73s
2. Wall time: The actual real world time spent in executing the task.
- For Model with Shared features:

▪ Avg Wall time over 3 trails: 24.9s

• For Model with Unshared features:

▪ Avg Wall time over 3 trails: 40.8s

These parameters supports the claim that the Actor critic model with shared features is better than the model with shared features.

In []:

1

In []:

1