


```
In [8]: 1 import numpy as np
        2 from enum import Enum
        3 import copy
```

Consider a standard grid world, where only 4 (up, down, left, right) actions are allowed and the agent deterministically moves accordingly, represented as below. Here yellow is the start state and white is the goal state.

Say, we define our MDP as:

- S: 121 (11 x 11) cells
- A: 4 actions (up, down, left, right)
- P: Deterministic transition probability
- R: -1 at every step
- gamma: 0.9

Our goal is to find an optimal policy (shown in right).



```
In [9]: 1 # Above grid is defined as below:
        2 #   - 0 denotes an navigable tile
        3 #   - 1 denotes an obstruction/wall
        4 #   - 2 denotes the start state
        5 #   - 3 denotes an goal state
        6
        7 # Note: Here the upper left corner is defined as (0, 0)
        8 #       and lower right corner as (m-1, n-1)
        9
       10 # Optimal Path: RIGHT RIGHT UP UP LEFT LEFT UP UP UP UP UP UP LEFT LEFT DOWN DOWN LEFT LEFT
       11
       12
       13 GRID_WORLD = np.array([
       14     [1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1],
       15     [1, 0, 0, 0, 0, 0, 0, 0, 1, 0, 1],
       16     [1, 1, 1, 0, 1, 0, 1, 1, 1, 0, 1],
       17     [1, 3, 0, 0, 1, 0, 1, 0, 1, 0, 1],
       18     [1, 1, 1, 1, 1, 0, 1, 0, 1, 0, 1],
       19     [1, 0, 0, 0, 1, 0, 1, 0, 0, 0, 1],
       20     [1, 1, 1, 0, 1, 0, 1, 0, 1, 1, 1],
       21     [1, 0, 0, 0, 0, 0, 0, 0, 1, 0, 1],
       22     [1, 1, 1, 1, 1, 1, 1, 0, 1, 0, 1],
       23     [1, 0, 0, 0, 0, 2, 0, 0, 0, 0, 1],
       24     [1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1]
       25 ])
```

Actions

```
In [10]: 1 class Actions(Enum):
2     UP     = (0, (-1, 0)) # index = 0, (xaxis_move = -1 and yaxis_move = 0)
3     DOWN   = (1, (1, 0))  # index = 1, (xaxis_move = 1 and yaxis_move = 0)
4     LEFT    = (2, (0, -1)) # index = 2, (xaxis_move = 0 and yaxis_move = -1)
5     RIGHT   = (3, (0, 1))  # index = 3, (xaxis_move = 0 and yaxis_move = 1)
6
7     def get_action_dir(self):
8         _, direction = self.value
9         return direction
10
11     @property
12     def index(self):
13         indx, _ = self.value
14         return indx
15
16     @classmethod
17     def from_index(cls, index):
18         action_index_map = {a.index: a for a in cls}
19         return action_index_map[index]
```

```
In [11]: 1 # How to use Action enum
2 for a in Actions:
3     print(f"name: {a.name}, action_id: {a.index}, direction_to_move: {a.get_action_dir()}")
4
5 print("\n-----\n")
6
7 # find action enum from index 0
8 a = Actions.from_index(3)
9 print(f"3 index action is: {a.name}")
```

name: UP, action\_id: 0, direction\_to\_move: (-1, 0)  
name: DOWN, action\_id: 1, direction\_to\_move: (1, 0)  
name: LEFT, action\_id: 2, direction\_to\_move: (0, -1)  
name: RIGHT, action\_id: 3, direction\_to\_move: (0, 1)

-----

3 index action is: RIGHT

Policy

```
In [12]: 1 class BasePolicy:
2     def update(self, *args):
3         pass
4
5     def select_action(self, state_id: int) -> int:
6         raise NotImplemented
7
8
9 class DeterministicPolicy(BasePolicy):
10     def __init__(self, actions: np.ndarray):
11         # actions: its a 1d array (|S| size) which contains action for each state
12         self.actions = actions
13
14     def update(self, state_id, action_id):
15         assert state_id < len(self.actions), f"Invalid state_id {state_id}"
16         assert action_id < len(Actions), f"Invalid action_id {action_id}"
17         self.actions[state_id] = action_id
18
19     def select_action(self, state_id: int) -> int:
20         assert state_id < len(self.actions), f"Invalid state_id {state_id}"
21         return self.actions[state_id]
```

## Environment

```
In [13]: 1 class Environment:
2     def __init__(self, grid):
3         self.grid = grid
4         m, n = grid.shape
5         self.num_states = m*n
6
7     def xy_to_posid(self, x: int, y: int):
8         _, n = self.grid.shape
9         return x*n + y
10
11    def posid_to_xy(self, posid: int):
12        _, n = self.grid.shape
13        return (posid // n, posid % n)
14
15    def isvalid_move(self, x: int, y: int):
16        m, n = self.grid.shape
17        return (x >= 0) and (y >= 0) and (x < m) and (y < n) and (self.grid[x, y] != 1)
18
19    def find_start_xy(self) -> int:
20        m, n = self.grid.shape
21        for x in range(m):
22            for y in range(n):
23                if self.grid[x, y] == 2:
24                    return (x, y)
25        raise Exception("Start position not found.")
26
27    def find_path(self, policy: BasePolicy) -> str:
28        max_steps = 50
29        steps = 0
30
31        P, R = self.get_transition_prob_and_expected_reward()
32        num_actions, num_states = R.shape
33        all_possible_state_posids = np.arange(num_states)
34
35        path = ""
36        curr_x, curr_y = self.find_start_xy()
37        while (self.grid[curr_x, curr_y] != 3) and (steps < max_steps):
38            curr_posid = self.xy_to_posid(curr_x, curr_y)
39            action_id = policy.select_action(curr_posid)
40            next_posid = np.random.choice(
41                all_possible_state_posids, p=P[action_id, curr_posid])
42            action = Actions.from_index(action_id)
43            path += f" {action.name}"
44            curr_x, curr_y = self.posid_to_xy(next_posid)
45            steps += 1
46        return path
47
48    def get_transition_prob_and_expected_reward(self): #  $P(s_{next} | s, a)$ ,  $R(s, a)$ 
49        m, n = self.grid.shape
50        num_states = m*n
51        num_actions = len(Actions)
52        P = np.zeros((num_actions, num_states, num_states))
53        R = np.zeros((num_actions, num_states))
54        for a in Actions:
55            for x in range(m):
56                for y in range(n):
57                    xmove_dir, ymove_dir = a.get_action_dir()
58                    xnew, ynew = x + xmove_dir, y + ymove_dir # find the new co-ordinate after the action a
59
60                    posid = self.xy_to_posid(x, y)
61                    new_posid = self.xy_to_posid(xnew, ynew)
62
63
64                    if self.grid[x, y] == 3:
65                        # the current state is a goal state
66                        P[a.index, posid, posid] = 1
67                        R[a.index, posid] = 0
68                    elif (self.grid[x, y] == 1) or (not self.isvalid_move(xnew, ynew)):
69                        # the current state is a block state or the next state is invalid
70                        P[a.index, posid, posid] = 1
71                        R[a.index, posid] = -1
72                    else:
73                        # action a is valid and goes to a new position
74                        P[a.index, posid, new_posid] = 1
75                        R[a.index, posid] = -1
76        return P, R
```

## Policy Iteration

```

In [14]: 1 def policy_evaluation(P: np.ndarray, R: np.ndarray, gamma: float,
2         policy: BasePolicy, theta: float,
3         init_V: np.ndarray=None):
4     num_actions, num_states = R.shape
5
6     # Please try different starting point for V you will find it will always
7     # converge to the same V_pi value.
8     if init_V is None:
9         init_V = -2.5*np.ones(num_states) # Different Starting point for V
10    V = copy.deepcopy(init_V)
11
12    delta = 100.0
13    while delta > theta:
14        delta = 0.0
15        for state_id in range(num_states):
16            action_id = policy.select_action(state_id)
17            v_old = V[state_id]
18            # Following equation is a different way of writing the same equation given in the slide.
19            # Note here R is an expected reward term.
20            V[state_id] = R[action_id, state_id] + gamma * np.dot(P[action_id, state_id], V)
21            delta = max(delta, abs(V[state_id] - v_old))
22    return V
23
24
25 def policy_improvement(P: np.ndarray, R: np.ndarray, gamma: float,
26                        policy: BasePolicy, V: np.ndarray):
27     num_actions, num_states = R.shape
28     policy_stable = True
29     for state_id in range(num_states):
30         old_action_id = policy.select_action(state_id)
31
32         # your code here
33         r = R[:, state_id]
34         p = P[:, state_id]
35         new_action_id = np.argmax(r + gamma * np.dot(p, V)) # update new_action_id based on the value function.
36
37         policy.update(state_id, new_action_id)
38         if old_action_id != new_action_id:
39             policy_stable = False
40     return policy_stable
41
42
43 def policy_iteration(P: np.ndarray, R: np.ndarray, gamma: float,
44                     theta: float=1e-3, init_policy: BasePolicy = None):
45     num_actions, num_states = R.shape
46
47     # Please try exploring different policies you will find it will always
48     # converge to the same optimal policy for valid states.
49     if init_policy is None:
50         # Say initial policy = all down actions.
51         init_policy = DeterministicPolicy(actions=np.ones(num_states, dtype=int)) # Changed All initial actions to go
52
53     # creating a copy of a initial policy
54     policy = copy.deepcopy(init_policy)
55     policy_stable = False
56     while not policy_stable:
57         V = policy_evaluation(P, R, gamma, policy, theta)
58         policy_stable = policy_improvement(P, R, gamma, policy, V)
59     return policy, V

```

Value Iteration

```
In [15]: 1 # Directly find the optimal value function
2 def get_optimal_value(P: np.ndarray, R: np.ndarray, gamma: float,
3                     theta: float, init_V: np.ndarray=None):
4     num_actions, num_states = R.shape
5
6     # Please try different starting point for V you will find it will always
7     # converge to the same V_star value.
8     if init_V is None:
9         init_V = -6.75* np.ones(num_states) # Different Starting point for V
10    V = copy.deepcopy(init_V)
11
12    delta = 100.0
13    while delta > theta:
14        delta = 0.0
15        for state_id in range(num_states):
16            v_old = V[state_id]
17            q_sa = np.zeros(num_actions)
18            for a in Actions:
19                q_sa[a.index] = R[a.index, state_id] + gamma * np.dot(P[a.index, state_id], V)
20            V[state_id] = np.max(q_sa)
21            delta = max(delta, abs(V[state_id] - v_old))
22    return V
23
24
25 def value_iteration(P: np.ndarray, R: np.ndarray, gamma: float,
26                   theta: float=1e-3, init_V: np.ndarray=None):
27     V_star = get_optimal_value(P, R, gamma, theta, init_V)
28
29     num_actions, num_states = R.shape
30     policy = DeterministicPolicy(actions=np.zeros(num_states, dtype=int))
31     for state_id in range(num_states):
32         # Your code here
33         r = R[:, state_id]
34         p = P[:, state_id]
35         action_id = np.argmax(r + gamma * np.dot(p, V_star)) # update the action_id based on V_star
36
37         policy.update(state_id, action_id)
38
39     return policy, V_star
```

Experiments

```
In [16]: 1 def is_same_optimal_value(V1, V2, diff_theta=1e-3):
2         diff = np.abs(V1 - V2)
3         return np.all(diff < diff_theta)
```

```
In [17]: 1 seed = 0
2 np.random.seed(seed)
3
4 gamma = 0.9
5 theta = 1e-5
```

```
In [18]: 1 env = Environment(GRID_WORLD)
2 P, R = env.get_transition_prob_and_expected_reward()
```

Exp 1: Using Policy iteration algorithm find the optimal path from start to goal position

```
In [19]: 1 # # Start with random choice of init_policy.
2 # One such choice could be: init_policy = np.ones(env.num_states, dtype=int)
3 # Start with your own choice of init_policy
4 init_policy = DeterministicPolicy(actions=2 * np.ones(env.num_states, dtype=int)) # Intial actions set to 2: LEFT
5
6 pitr_policy, pitr_V_star = policy_iteration(P, R, gamma, theta=theta, init_policy=init_policy)
7 pitr_path = env.find_path(pitr_policy)
8 print(pitr_path)

RIGHT RIGHT UP UP LEFT LEFT UP UP UP UP UP UP LEFT LEFT DOWN DOWN LEFT LEFT
```

Exp 2: Using value iteration algorithm find the optimal path from start to goal position

```
In [20]: 1 vitr_policy, vitr_V_star = value_iteration(P, R, gamma, theta=theta)
2 vitr_path = env.find_path(vitr_policy)
3 print(vitr_path)
```

RIGHT RIGHT UP UP LEFT LEFT UP UP UP UP UP UP LEFT LEFT DOWN DOWN LEFT LEFT

Exp 3: Compare the optimal value function of policy iteration and value iteration algorithm

```
In [22]: 1 is_same_optimal_value(pitr_V_star, vitr_V_star)
```

Out[22]: True

Exp 4: Using initial guess for V as random values, find the optimal value function using policy evaluation and compare it with the optimal value function

```
In [23]: 1 # Start with random choice of init_V.
2 # One such choice could be: init_V = np.random.randn(env.num_states)
3 # Another choice could be: init_V = 10*np.ones(env.num_states)
4 # Start with your own choice of init_V
5 init_V = -169*np.random.rand(env.num_states) # Different Starting point for V
6
7 V_star = policy_evaluation(P, R, gamma, pitr_policy, theta, init_V)
8 is_same_optimal_value(pitr_V_star, V_star)
```

Out[23]: True

Exp 5: Using initial guess for V as random values, find the optimal value function using get\_optimal\_value and compare it with the optimal value function

```
In [24]: 1 # Start with random choice.
2 # One such choice could be: init_V = np.random.randn(env.num_states)
3 # Another choice could be: init_V = 10*np.ones(env.num_states)
4 # Start with your own choice of init_V
5 init_V = -1e7*np.random.random(env.num_states) # Different Starting point for V
6
7 V_star = get_optimal_value(P, R, gamma, theta, init_V)
8 is_same_optimal_value(vitr_V_star, V_star)
```

Out[24]: True

Exp Optional: Try changing the grid by adding multiple paths to the goal state and check if our policy\_iteration or value\_iteration algorithm is able to find optimal path. Redo the above experiments.

- 1 way to add another path would be GRID\_WORLD[4, 1] = 0

```
In [ ]: 1
```

```
In [25]: 1 from google.colab import drive
2 drive.mount('/content/drive')
```

Mounted at /content/drive

```
In [28]: 1 !jupyter nbconvert --to pdf /content/drive/MyDrive/DA6400_Value_and_Policy_Iteration_DA24S018.ipynb
```

[NbConvertApp] Converting notebook /content/drive/MyDrive/DA6400\_Value\_and\_Policy\_Iteration\_DA24S018.ipynb to pdf  
[NbConvertApp] Writing 63602 bytes to notebook.tex  
[NbConvertApp] Building PDF  
[NbConvertApp] Running xelatex 3 times: ['xelatex', 'notebook.tex', '-quiet']  
[NbConvertApp] Running bibtex 1 time: ['bibtex', 'notebook']  
[NbConvertApp] WARNING | bibtex had problems, most likely because there were no citations  
[NbConvertApp] PDF successfully created  
[NbConvertApp] Writing 68870 bytes to /content/drive/MyDrive/DA6400\_Value\_and\_Policy\_Iteration\_DA24S018.pdf

```
In [ ]: 1
```